

# Docker

## Core Concepts

- **Docker:** A platform for building, shipping, and running applications using containerization. It simplifies software installation, distribution, and management.
- **Containers:** Isolated runtime environments that package an application and its dependencies. They are a lightweight alternative to virtual machines (VMs).
- **Images:** Read-only templates used to create containers. They contain the application, libraries, and dependencies.
- **Docker Hub:** A public registry for storing and sharing Docker images.
- **Daemon:** A background process or service that runs continuously.

## Key Benefits

- **Isolation:** Containers isolate applications from each other and the host system, preventing conflicts.
- **Portability:** Docker images can run on any system with Docker installed.
- **Efficiency:** Containers consume fewer resources than VMs.
- **Organization:** Docker helps manage dependencies and keeps the system clean.
- **Security:** Containers limit the scope of potential security breaches.

## Basic Docker Commands

- `docker run <image>`: Creates and starts a container from an image.
  - `-d` or `--detach`: Runs the container in the background (detached mode).
  - `--name <container_name>`: Assigns a name to the container.
  - `-i` or `--interactive`: Keeps STDIN open even if not attached.
  - `-t` or `--tty`: Allocates a pseudo-TTY.
  - `--link <container_name>:<alias>`: Creates a link to another container (legacy).
- `docker ps`: Lists running containers.
  - `-a`: Lists all containers (running and stopped).
- `docker logs <container_name>`: Displays the logs of a container.
  - `-f` or `--follow`: Follow log output in real-time.
- `docker stop <container_name>`: Stops a running container.
- `docker restart <container_name>`: Restarts a stopped container.
- `docker exec <container_name> <command>`: Executes a command inside a running container.
- `docker create <image>`: Creates a container but does not start it.
- `docker rename <old_name> <new_name>`: Renames a container.
- `docker inspect <container_name>`: Displays detailed information about a container.
- `docker help <command>`: Shows help information for a specific Docker command.

## Important Concepts & Techniques

- **PID Namespace:** Isolates process IDs within a container.
- **Container Naming:** Use meaningful names or IDs for container management.
- **Container State:** Containers can be in various states (created, running, stopped, etc.).
- **Environment Independence:** Minimize dependencies on the host environment.
- **Read-only Filesystems:** Improves security and prevents unintended modifications. Use the `--read-only` flag.
- **Volumes:** Used to persist data and share files between the host and containers. `-v <host_path>:<container_path>`
- **tmpfs:** Creates temporary, in-memory filesystems. `--tmpfs <container_path>`
- **Environment Variables:** Inject configuration data into containers. `-e <VARIABLE>=<value>`

## Troubleshooting

- **Port Conflicts:** Occur when multiple containers try to use the same port.
- **Dependency Order:** Ensure containers are started in the correct order based on dependencies.
- **Image Installation:** Verify that the image is installed correctly.

## Workflow Summary

1. **Pull/Build Image:** Obtain a Docker image from a registry or build your own.
2. **Create Container:** Create a container from the image using `docker run` or `docker create`.
3. **Start Container:** Start the container using `docker start`.
4. **Manage Container:** Monitor, stop, restart, or execute commands within the container.

## Networking

### Key Concepts

- **Protocols:** Languages for network communication (e.g., HTTP).
- **Network Interface:** A point of connection to a network, having an IP address.
- **Port:** A specific channel on an interface for a process (e.g., port 80 for HTTP).
- **Bridge:** Connects multiple networks.
- Docker abstracts host networks from containers.
- Containers get unique IP addresses within a Docker network.
- Networks are first-class entities, managed with `docker network` commands.

### Default Networks

- **bridge:** Default network, inter-container connectivity on a single host (Legacy, not recommended).
- **host:** No network isolation; container uses the host's network directly.
- **none:** No external network connectivity for the container.

### Network Scope

- local: Confined to a single host.
- global: Created on all cluster nodes, but doesn't route between them.
- swarm: Spans all hosts in a Docker Swarm cluster.

## User-Defined Bridge Networks

- Use `docker network create` to create custom bridge networks.
- Enable modern Docker features like service discovery and load balancing.

## Creating a Network

`docker network ls` # List networks

`docker network create <network_name>` # Create a network

## Bridge Networks (Single-Host)

- Default: `docker0` (virtual interface).
- Containers get private IPs, routed via Docker network.
- User-defined bridges: `docker network create --driver bridge ...`
  - `--attachable`: Allows attaching/detaching containers.
  - `--subnet, --ip-range`: Custom IP addressing.
- `docker network connect`: Attaches container to a network.

## Exploring Networks

- `ip addr`: Shows container network interfaces (`lo`, `eth0`, etc.).
- `nmap`: Scans network for connected devices.
- DNS-based service discovery: Containers discover each other by name (`hostname = container name`).

## Beyond Bridge Networks

- **Underlay Networks** (`macvlan`, `ipvlan`): First-class network addresses, routable from host network. Linux hosts only.
- **Overlay Networks (Swarm Mode)**: Multi-host aware bridge, inter-container connections across swarm nodes.

## Special Networks

- `--network host`: Container shares host's network namespace (no isolation). Access to localhost services.
- `--network none`: Container is network-isolated (only loopback interface).

## Inbound Traffic (NodePort Publishing)

- -p host\_port:container\_port: Forwards traffic from host to container.
- docker port container\_name: Lookup port mappings.

## Networking Caveats

- **No Firewalls:** Containers on same network have unrestricted access to each other. Application-level security is crucial.
- **Custom DNS:**
  - --hostname: Sets container hostname (resolves to container IP).
  - --dns: Specifies DNS servers for container.
  - --dns-search: Specifies DNS search domain (appends to hostnames).
  - --add-host: Adds custom hostname-to-IP mapping.

## Image Packaging

### Building from a Container

1. docker run: Create container from existing image.
2. Modify filesystem (install software, create files, etc.).
3. docker commit: Create new image from container's changes.

### Reviewing Changes

- docker container diff container\_name (A=added, C=changed, D=deleted).

### Committing Images

- docker container commit -a "author" -m "message" container\_name new\_image\_name
  - -a: Author.
  - -m: Commit message.
- Image attributes, such as environment variables, working directory, exposed ports, volumes, entrypoint, and command are all saved in the image.

### Union Filesystems (UFS)

- Images are built in layers.
- Changes are written to new layers.
- Reading a file: UFS searches layers top-down.
- File deletion: Adds a "delete" record to the top layer.
- File changes: Adds the changed file to the top layer.

## Image Distribution

# Key Concepts

- **Distribution Spectrum:** Ranging from simple (hosted public registries) to complex (custom distribution).
- **Selection Criteria:**
  - **Cost:** Free to expensive.
  - **Visibility:** Public vs. private.
  - **Transport Speed/Bandwidth:** Installation speed.
  - **Longevity Control:** Risk of third-party changes.
  - **Availability Control:** Ability to resolve availability issues.
  - **Access Control:** Protection from unauthorized modification.
  - **Artifact Integrity:** Ensuring image authenticity.
  - **Artifact Confidentiality:** Protecting sensitive content.
  - **Requisite Expertise:** Skill level needed.

## Distribution Methods

1. **Hosted Public Registries (e.g., Docker Hub):** Simple, high visibility. Limited control over longevity and availability.
2. **Hosted Private Registries:** Offer access control and confidentiality. Still rely on a third party for availability and longevity.
3. **Private Registries:** Run your own registry using registry software. Greater control over all aspects, but increased responsibility. Suitable for local, corporate, or private cloud networks.
4. **Custom Image Distribution:** Maximum flexibility. Requires building your own infrastructure (e.g., SFTP, HTTP downloads).
5. **Image Source Distribution:** Distribute Dockerfiles instead of images. Requires users to build images themselves. Suitable for open-source projects.

## Services

### Key Concepts

- **Service:** A discoverable and available process/functionality over a network. Abstraction simplifies management.
- **Task:** A unit of work within a service, often a container.
- **Swarm Mode:** Enables Docker service management features. Requires initialization (docker swarm init).
- **Orchestrator (e.g., Docker Swarm):** Automates service lifecycle, ensuring desired state is maintained. Tracks desired state vs. current state.
- **Replication:** Running multiple instances (replicas) of a service for availability.

## Service Lifecycle & Management

1. **Service Creation:** docker service create defines a service (name, image, port mapping).
2. **Resurrection:** Swarm automatically restarts failed containers to maintain the desired number of replicas.
3. **Scaling:** docker service scale adjusts the number of service replicas.

4. **Inspection:** docker service inspect shows the service's desired state definition.

5. **Replication Modes:**

- **Replicated:** Maintains a fixed number of replicas.
- **Global:** Runs one replica on each node in the swarm.

1. **Automated Rollout:** docker service update updates a service's configuration (e.g., image).

- `--update-order`: stop-first (stop old before starting new).
- `--update-parallelism`: Number of replicas updated simultaneously.
- `--update-delay`: Time between updates.

## Image Build Patterns

- **All-in-One:** Includes build tools and runtime dependencies in a single image (simple, large).
- **Build Plus Runtime:** Separate build image and a smaller runtime image (better security, smaller size).
- **Build Plus Multiple Runtimes:** Slim runtime image with variations for debugging or specialized use cases (multi-stage builds).

## Image Metadata

- **Labels:** Use LABEL instruction to add metadata (application name, version, build date, VCS commit).
- **Dockerfile & Manifests:** Include these in the image filesystem for traceability.
- **Orchestration with Make:** Use make to automate the build process (metadata gathering, artifact building, image building, testing, tagging).

## Image Testing

- **Dockerfile Linting:** Use tools like hadolint to check for best practices.
- **Container Structure Test (CST):** Verifies file permissions, command execution, metadata.
- **Vulnerability Scanning:** Scan images for known vulnerabilities before publishing.

## Image Tagging

- Tags are mutable, human-readable pointers to image IDs.
- **Continuous Delivery with Unique Tags:** Use unique build IDs for each image and promote them through stages.
- **Configuration Image per Deployment Stage:** Create a generic app image and separate config images for each environment.
- **Semantic Versioning:** Use Major.Minor.Patch for releases to communicate the level of change.

## Swarm

# Core Concepts

- **Swarm:** Docker's built-in clustering & orchestration. Enables deploying apps across multiple hosts for scalability & availability.
- **Nodes:** Machines in a Swarm cluster.
  - **Manager:** Orchestrates services, maintains cluster state. Requires a majority to be available.
  - **Worker:** Executes tasks (containers) as instructed by managers.
- **Services:** Define application processes. Swarm creates tasks to realize service definitions.
- **Tasks:** Containerized process scheduled & run once. Replaced if they fail based on restart policy.
- **Networks:** Overlay networks for service communication. Encrypted for security.
- **Volumes:** Persistent storage for service data. Local to a node by default.
- **Configs/Secrets:** Externalized configuration data. Secrets are handled securely.

## Deployment & Management

- **Swarm Initialization:** `docker swarm init` (on a manager node).
- **Node Joining:** `docker swarm join` (using join tokens).
- **Docker Stack:** A named collection of services, networks, and volumes. Defined using Docker Compose (YAML).
  - `docker stack deploy`: Creates/updates a stack.
  - `docker stack ps`: Lists tasks in a stack.
  - `docker stack rm`: Removes a stack.
- **Service Management:**
  - `docker service update`: Modifies service configuration.
  - `docker service ps`: Lists tasks for a service.
  - `docker service logs`: Shows logs for a service.
- **Config Management:**
  - `docker config create`: Creates a config.
  - `docker config inspect`: Inspects a config.
  - `docker config rm`: Removes a config.
- **Secret Management:**
  - `docker secret create`: Creates a secret.
  - `docker secret inspect`: Inspects a secret.
  - `docker secret rm`: Removes a secret.

## Compose File (`docker-compose.yml`)

- **version:** Specifies Compose file version.
- **services:** Defines services.
  - **image:** Container image to use.
  - **ports:** Port mappings (host:container).
  - **networks:** Networks to attach the service to.
  - **volumes:** Volume mounts (host\_path:container\_path).
  - **environment:** Environment variables.
  - **secrets:** Secrets to mount.
  - **configs:** Configs to mount.

- **deploy:** Deployment configuration.
  - **replicas:** Number of task replicas.
  - **restart\_policy:** Defines restart behavior.
  - **update\_config:** Configures update strategy.
  - **resources:** Resource limits (CPU, memory).
- **networks:** Defines networks.
  - **driver:** Network driver (e.g., overlay).
  - **driver\_opts:** Driver options (e.g., encryption).
  - **attachable:** Allows standalone containers to attach to this network.
- **volumes:** Defines volumes.
- **secrets:** Defines external secrets.
  - **external: true:** Indicates secret is managed outside the Compose file.
- **configs:** Defines external configs.
  - **external: true:** Indicates config is managed outside the Compose file.

## Key Considerations

- **Service Discovery:** Services discoverable by name within a Docker network.
- **Load Balancing:** Swarm provides built-in load balancing across service replicas.
- **Security:** Use encrypted networks. Secure published ports with TLS.
- **Secrets Management:** Avoid environment variables for secrets. Use Docker secrets mounted as files.
- **Immutable Configs:** Configs are immutable. Create new configs on change.
- **Restart Policies:** Define robust restart strategies.
- **Update Configuration:** Control update rollout (parallelism, delay).
- **Resource Limits:** Set CPU and memory limits for services.
- **Dependencies:** Use `depends_on` to define service startup order.
- **Rolling Updates:** Swarm updates services in batches, stopping old tasks before starting new ones.

## Core Concepts (Redux)

- **Nodes:** Machines in the Swarm cluster (managers & workers). `docker node ls` shows node status.
- **Services:** Application definitions (e.g., API, database). Defined in `docker-compose.yml`.
- **Tasks:** Individual container instances running as part of a service.
- **Stacks:** Collection of related services defined in a `docker-compose.yml` file. Deployed with `docker stack deploy`.
- **Secrets:** Secure way to manage sensitive data (passwords, keys). Created with `docker secret create`.
- **Configs:** Non-sensitive configuration data.
- **Networks:** Overlay networks enable communication between services.

## Key Commands (Redux)

- `docker node ls`: List nodes in the Swarm.
- `docker stack deploy --compose-file docker-compose.yml <stack_name>`: Deploy a stack.
- `docker service ls`: List services.



- `docker service logs <service_name>`: View service logs.
- `docker service ps <service_name>`: List tasks for a service.
- `docker secret create <secret_name> <file>`: Create a secret.
- `docker stack rm <stack_name>`: Remove a stack.

## Deployment Process (Redux)

1. Initialize a Swarm (on a manager node).
2. Define services in a `docker-compose.yml` file, specifying images, ports, networks, secrets, and configs.
3. Create secrets (if needed).
4. Deploy the stack using `docker stack deploy`.
5. Verify service status with `docker service ls` and `docker service ps`.
6. Access the application through published ports on any node in the Swarm.

## Networking (Redux)

- **Ingress Network**: Special overlay network for external access to services.
- **Routing Mesh**: Exposes service ports on all nodes, routing traffic to available tasks.
- **Service Discovery**: Services use DNS to find each other by name.
- **Overlay Networks**: Isolate communication between services.

## Task Placement (Redux)

- **Replicated Mode**: Swarm maintains the desired number of replicas.
- **Global Mode**: One task per node.
- **Placement Constraints**: Control where tasks run using node attributes (e.g., `node.role == worker`, `node.labels.zone == public`). Use `docker service update --constraint-add`.

## Resource Management (Redux)

- **Limits**: Maximum resources a container can use (e.g., CPU, memory).
- **Reservations**: Guaranteed minimum resources.
- Use `deploy.resources.limits` and `deploy.resources.reservations` in `docker-compose.yml`.

## Important Notes (Redux)

- Swarm requires all cluster-level resources (networks, secrets, configs) to exist before deploying services.
- Service tasks are ephemeral; updates replace containers.
- Data in default volumes is local to a node. Use volume plugins for shared storage.
- No firewalls between containers on a Docker overlay network.
- Application-level authentication is crucial for service-to-service communication.
- Swarm load balances connections, not HTTP requests, by default.